AI Agent Control Plane: Agent Guidance and Change Lifecycle

Feedforward

Feedback

Time ↓

Feedforward

runtime entrypoints / context routing →

instructions / skills →

xecution policies role charters →

MCP declarations / provider adapters →

ontracts / schema →

hooks / provenance →

Token Economics review→

spec-driven RPI check →

Humans and/or agents promote durable lessons into instructions, skills, hooks, policies, role charters, MCP declarations, contracts, schemas, docs, adapters, or evaluation targets

Coding Agent

↻ Load governed context

↻ Propose artifact / diff

↻ Self-correct from feedback

Initial generation

proposed artifact / diff

← git diff review

← dependency / boundary check

← validation coverage

← pattern / policy scan

← lint / static check

Self-correction Loop

Human review

durable artifact proposal

human

human

Promotion

Provenance / load report / execution evidence observation → proposed change → review → promoted asset

Behavior Steering

Architectural Fitness

Maintainability

Governance and Validation Pipeline

← Specification / control artifact

← git diff review

← architecture / ADR review

← hook shell syntax

← source alignment review

←instruction routing review

← artifact ownership check

←workspace MCP config check

← downstream project validation

blue-gray = computational

red = inferential

Adapted from Martin Fowler's Harness Engineering model